


Qualitative examples of label inaccuracies in the AVA-Speech dataset




*The red colored boundary region shows inaccurate labels

Filename: 5BDj0ow5hnA.wav

chunk: 632799 to 716639

Label: SPEECH_WITH_NOISE

Our observation: There are noise-only segments as well as speech-with-noise segments




*The red colored boundary region shows inaccurate labels

Filename: 5BDj0ow5hnA.wav

chunk: 1395200 to 1474399

Label: CLEAN_SPEECH

Our observation: There is a silence segment of ~0.4 sec, starting from 3.28 sec




*The red colored boundary region shows inaccurate labels

Filename: 5BDj0ow5hnA.wav

chunk: 5526880 to 5564639

Label: CLEAN_SPEECH



Our observation: There is a silence segment of ~ 0.8 sec, starting from 0.64 sec



*The red colored boundary region shows inaccurate labels

Filename: 9bK05eBt1GM.wav
chunk: 13994880 to 14020480
Label: CLEAN_SPEECH

Our observation: This segment seems to be a noise segment (music) rather than a speech segment





*The red colored boundary region shows inaccurate labels

Filename: 9bK05eBt1GM.wav

chunk: 25440 to 94079

Label: SPEECH_WITH_NOISE

Our observation: There are noise-only segments as well as speech-with-noise segments



*The red colored boundary region shows inaccurate labels

Filename: 9bK05eBt1GM.wav

chunk: 4756480 to 4989279

Label: CLEAN_SPEECH

Our observation: There are several silent segments in addition to speech segments